

# FINAL COURSE PROJECT: TITANIC SURVIVAL DATASET ANALYSIS

In this final project we are going to explore and study the Titanic survival dataset, and structure all the analysis information in a PDF report. These are the sections that we will go through:

- Brief description of the data set and a summary of its attributes
- Initial plan for data exploration
- Actions taken for data cleaning and feature engineering
- Key Findings and Insights, which synthesizes the results of Exploratory Data Analysis in an insightful and actionable manner
- Formulating at least 3 hypothesis about this data
- Conducting a formal significance test for one of the hypotheses and discuss the results
- Suggestions for next steps in analyzing this data
- Quality of this data set and a request for additional data if needed

The dataset can be downloaded ([HERE](#))

## STEP 1: Dataset download, description and attributes summary

The sinking of the Titanic is one of the most infamous shipwrecks in history.

On April 15, 1912, during her maiden voyage, the widely considered "unsinkable" RMS Titanic sank after colliding with an iceberg. Unfortunately, there weren't enough lifeboats for everyone onboard, resulting in the death of 1502 out of 2224 passengers and crew.

While there was some element of luck involved in surviving, it seems some groups of people were more likely to survive than others.

Train.csv will contain the details of a subset of the passengers on board (891 to be exact) and importantly, will reveal whether they survived or not, also known as the "ground truth".

```
In [ ]: import kaggle

competition_name = 'titanic'
kaggle.api.competition_download_files(competition_name, path='titanic_data', quiet=

Warning: Your Kaggle API key is readable by other users on this system! To fix thi
s, you can run 'chmod 600 /home/cmadaria/.kaggle/kaggle.json'
Downloading titanic.zip to titanic_data
100% |██████████| 34.1k/34.1k [00:00<00:00, 4.68MB/s]
```

```
In [5]: import zipfile

# Define the path to the downloaded ZIP file
zip_file_path = 'titanic_data/titanic.zip'

# Unzip the file
with zipfile.ZipFile(zip_file_path, 'r') as zip_ref:
    zip_ref.extractall('titanic_data')
```

```
In [4]: import pandas as pd
```

```
In [5]: df = pd.read_csv('titanic_data/train.csv')
```

```
In [3]: # Display the first few rows of the dataset
print("First few rows of the dataset:")
print(df.head())
```

First few rows of the dataset:

	PassengerId	Survived	Pclass	\
0	1	0	3	
1	2	1	1	
2	3	1	3	
3	4	1	1	
4	5	0	3	

	Name	Sex	Age	SibSp	\
0	Braund, Mr. Owen Harris	male	22.0	1	
1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	
2	Heikkinen, Miss. Laina	female	26.0	0	
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	
4	Allen, Mr. William Henry	male	35.0	0	

	Parch	Ticket	Fare	Cabin	Embarked
0	0	A/5 21171	7.2500	NaN	S
1	0	PC 17599	71.2833	C85	C
2	0	STON/O2. 3101282	7.9250	NaN	S
3	0	113803	53.1000	C123	S
4	0	373450	8.0500	NaN	S

```
In [4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
#   Column      Non-Null Count  Dtype
---  -
0   PassengerId  891 non-null    int64
1   Survived     891 non-null    int64
2   Pclass       891 non-null    int64
3   Name         891 non-null    object
4   Sex          891 non-null    object
5   Age          714 non-null    float64
6   SibSp        891 non-null    int64
7   Parch        891 non-null    int64
8   Ticket       891 non-null    object
9   Fare         891 non-null    float64
10  Cabin        204 non-null    object
11  Embarked     889 non-null    object
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB
```

```
In [5]: df.shape[0]
```

```
Out[5]: 891
```

```
In [6]: df.columns.to_list()
```

```
Out[6]: ['PassengerId',
         'Survived',
         'Pclass',
         'Name',
         'Sex',
         'Age',
         'SibSp',
         'Parch',
         'Ticket',
         'Fare',
         'Cabin',
         'Embarked']
```

```
In [7]: df.describe()
```

```
Out[7]:
```

	PassengerId	Survived	Pclass	Age	SibSp	Parch	Fare
<b>count</b>	891.000000	891.000000	891.000000	714.000000	891.000000	891.000000	891.000000
<b>mean</b>	446.000000	0.383838	2.308642	29.699118	0.523008	0.381594	32.204208
<b>std</b>	257.353842	0.486592	0.836071	14.526497	1.102743	0.806057	49.693429
<b>min</b>	1.000000	0.000000	1.000000	0.420000	0.000000	0.000000	0.000000
<b>25%</b>	223.500000	0.000000	2.000000	20.125000	0.000000	0.000000	7.910400
<b>50%</b>	446.000000	0.000000	3.000000	28.000000	0.000000	0.000000	14.454200
<b>75%</b>	668.500000	1.000000	3.000000	38.000000	1.000000	0.000000	31.000000
<b>max</b>	891.000000	1.000000	3.000000	80.000000	8.000000	6.000000	512.329200

```
In [8]: df.dtypes
```

```
Out[8]: PassengerId      int64
Survived      int64
Pclass      int64
Name      object
Sex      object
Age      float64
SibSp      int64
Parch      int64
Ticket      object
Fare      float64
Cabin      object
Embarked      object
dtype: object
```

```
In [9]: print(df.describe(include=['object', 'category']))
```

```

              Name  Sex  Ticket      Cabin  Embarked
count          891   891      891        204        889
unique          891     2       681        147         3
top  Dooley, Mr. Patrick  male   1601  B96 B98         S
freq           1    577         7         4        644
```

```
In [10]: df['Embarked'].value_counts()
```

```
Out[10]: Embarked
S      644
C      168
Q       77
Name: count, dtype: int64
```

```
In [11]: df['Survived'].value_counts()
```

```
Out[11]: Survived
0      549
1      342
Name: count, dtype: int64
```

```
In [13]: means = df.describe().loc['mean']
medians = df.describe().loc['50%']
quantile_25 = df.describe().loc['25%']
quantile_75 = df.describe().loc['75%']
ranges = df.describe().loc['max'] - df.describe().loc['min']
print(means, medians, quantile_25, quantile_75, ranges)
```

```
PassengerId    446.000000
Survived        0.383838
Pclass         2.308642
Age            29.699118
SibSp          0.523008
Parch          0.381594
Fare           32.204208
Name: mean, dtype: float64 PassengerId    446.0000
Survived        0.0000
Pclass          3.0000
Age             28.0000
SibSp           0.0000
Parch           0.0000
Fare            14.4542
Name: 50%, dtype: float64 PassengerId    223.5000
Survived        0.0000
Pclass          2.0000
Age             20.1250
SibSp           0.0000
Parch           0.0000
Fare             7.9104
Name: 25%, dtype: float64 PassengerId    668.5
Survived        1.0
Pclass          3.0
Age             38.0
SibSp           1.0
Parch           0.0
Fare            31.0
Name: 75%, dtype: float64 PassengerId    890.0000
Survived        1.0000
Pclass          2.0000
Age             79.5800
SibSp           8.0000
Parch           6.0000
Fare           512.3292
dtype: float64
```

```
In [14]: stats_df = df.describe()
stats_df.loc['range'] = stats_df.loc['max'] - stats_df.loc['min']

out_fields = ['mean', '25%', '50%', '75%', 'range']
stats_df = stats_df.loc[out_fields]
stats_df.rename({'50%': 'median'}, inplace=True)
stats_df
```

Out[14]:

	PassengerId	Survived	Pclass	Age	SibSp	Parch	Fare
<b>mean</b>	446.0	0.383838	2.308642	29.699118	0.523008	0.381594	32.204208
<b>25%</b>	223.5	0.000000	2.000000	20.125000	0.000000	0.000000	7.910400
<b>median</b>	446.0	0.000000	3.000000	28.000000	0.000000	0.000000	14.454200
<b>75%</b>	668.5	1.000000	3.000000	38.000000	1.000000	0.000000	31.000000
<b>range</b>	890.0	1.000000	2.000000	79.580000	8.000000	6.000000	512.329200

Attribute	Data Type	Description
<b>PassengerId</b>	Numeric	A unique identifier for each passenger.
<b>Survived</b>	Binary (0/1)	Indicates whether the passenger survived (1 = survived, 0 = did not survive).
<b>Pclass</b>	Categorical	Passenger class (1 = First, 2 = Second, 3 = Third).
<b>Name</b>	Text	Full name of the passenger, often including titles (e.g., Mr., Mrs.).
<b>Sex</b>	Categorical	Gender of the passenger (male or female).
<b>Age</b>	Numeric	Age of the passenger in years. Some values are missing and need imputation.
<b>SibSp</b>	Numeric	Number of siblings and/or spouses aboard the Titanic.
<b>Parch</b>	Numeric	Number of parents and/or children aboard the Titanic.
<b>Ticket</b>	Text	Ticket number, which can include letters and numbers.
<b>Fare</b>	Numeric	Fare paid for the ticket.
<b>Cabin</b>	Text	Cabin number. Many values are missing.
<b>Embarked</b>	Categorical	Port of embarkation: C = Cherbourg, Q = Queenstown, S = Southampton.

## STEP 2: Initial plan for Data Exploration

These will be the performed steps in the Data Exploration stage:

- 1- General Data Overview
- 2- Univariate analysis
- 3- Correlation analysis
- 4- Log transformation
- 5- Handling Duplicates
- 6- Missing values
- 7- Outliers
- 8- Scaling
- 9- New variables creation
- 10- Encoding categorical features
- 11- Feature Selection

## STEP 3: Actions taken for data cleaning and feature engineering

```
In [8]: import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
```

# 1- General Data Overview

In [15]: `df.head()`

Out[15]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500	NaN
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...)	female	38.0	1	0	PC 17599	71.2833	C85
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250	NaN
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000	C123
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500	NaN

In [ ]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
#   Column          Non-Null Count  Dtype
---  -
0   PassengerId     891 non-null   int64
1   Survived        891 non-null   int64
2   Pclass          891 non-null   int64
3   Name            891 non-null   object
4   Sex             891 non-null   object
5   Age             714 non-null   float64
6   SibSp           891 non-null   int64
7   Parch           891 non-null   int64
8   Ticket          891 non-null   object
9   Fare            891 non-null   float64
10  Cabin           204 non-null   object
11  Embarked        889 non-null   object
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB
```

In [ ]: `df.shape[0]`

891

In [ ]: `df.columns.to_list()`

```
[ 'PassengerId',
  'Survived',
  'Pclass',
  'Name',
  'Sex',
  'Age',
  'SibSp',
  'Parch',
  'Ticket',
  'Fare',
  'Cabin',
  'Embarked' ]
```

```
In [ ]: df.describe()
```

	PassengerId	Survived	Pclass	Age	SibSp	Parch	Fare
<b>count</b>	891.000000	891.000000	891.000000	714.000000	891.000000	891.000000	891.000000
<b>mean</b>	446.000000	0.383838	2.308642	29.699118	0.523008	0.381594	32.204208
<b>std</b>	257.353842	0.486592	0.836071	14.526497	1.102743	0.806057	49.693429
<b>min</b>	1.000000	0.000000	1.000000	0.420000	0.000000	0.000000	0.000000
<b>25%</b>	223.500000	0.000000	2.000000	20.125000	0.000000	0.000000	7.910400
<b>50%</b>	446.000000	0.000000	3.000000	28.000000	0.000000	0.000000	14.454200
<b>75%</b>	668.500000	1.000000	3.000000	38.000000	1.000000	0.000000	31.000000
<b>max</b>	891.000000	1.000000	3.000000	80.000000	8.000000	6.000000	512.329200

```
In [ ]: df.dtypes
```

```
PassengerId    int64
Survived        int64
Pclass          int64
Name            object
Sex             object
Age            float64
SibSp           int64
Parch           int64
Ticket          object
Fare            float64
Cabin           object
Embarked        object
dtype: object
```

```
In [ ]: print(df.describe(include=['object', 'category']))
```

```

           Name  Sex  Ticket  Cabin  Embarked
count          891   891     891    204      889
unique          891     2     681    147       3
top  Dooley, Mr. Patrick  male  1601  B96 B98       S
freq             1   577       7      4     644
```

```
In [17]: df['Cabin'].value_counts()
```

```
Out[17]: Cabin
B96 B98      4
G6         4
C23 C25 C27  4
F2         3
C22 C26     3
..
C106        1
A19         1
D7          1
C118        1
E50         1
Name: count, Length: 147, dtype: int64
```

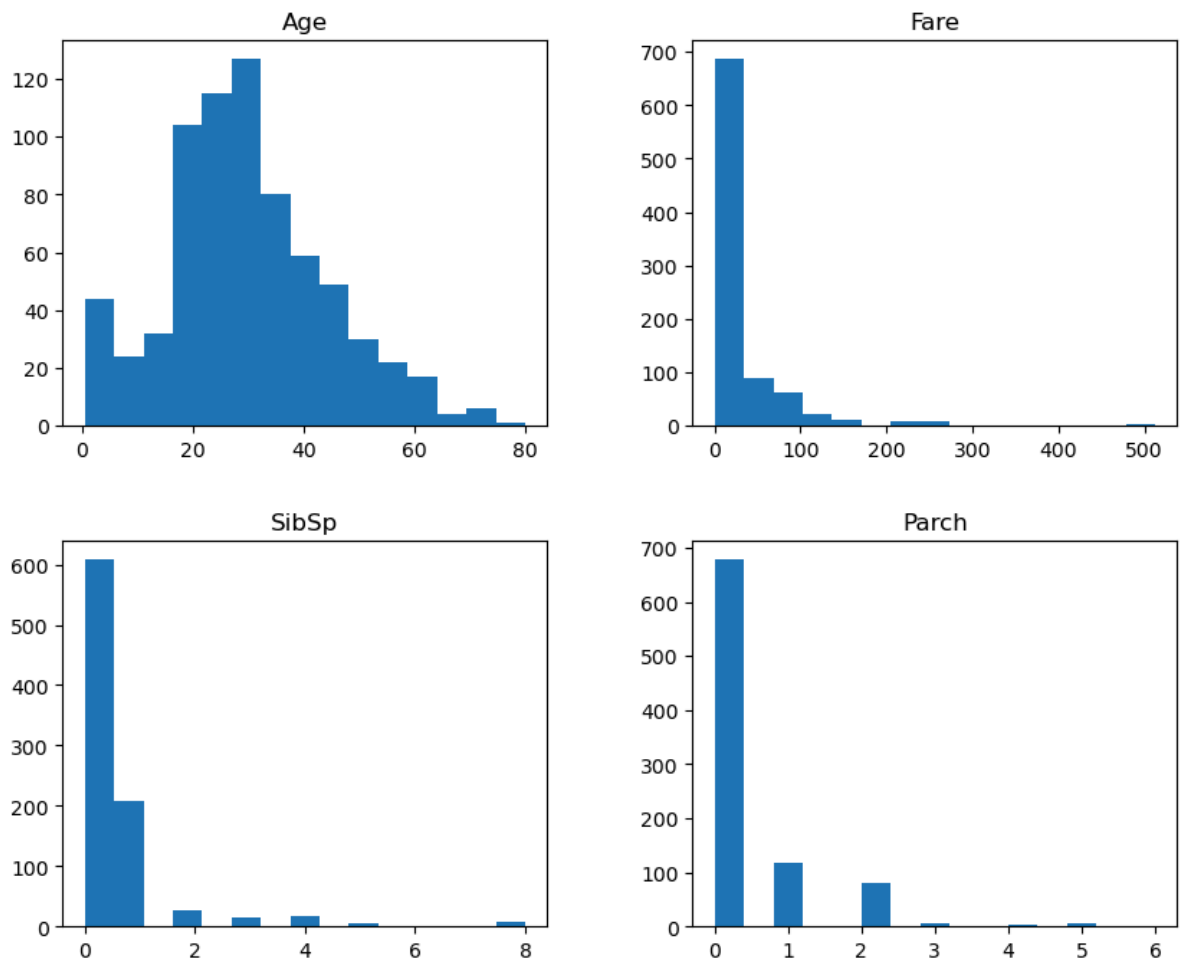
```
In [20]: df['Embarked'].value_counts()
```

```
Out[20]: Embarked
S      644
C      168
Q       77
Name: count, dtype: int64
```

## 2- Univariate Analysis



```
In [30]: # Plot histograms for numeric columns
numeric_cols = ['Age', 'Fare', 'SibSp', 'Parch']
df[numeric_cols].hist(bins=15, figsize=(10, 8), grid=False)
plt.suptitle("Histograms for Numeric Columns")
plt.show()
```

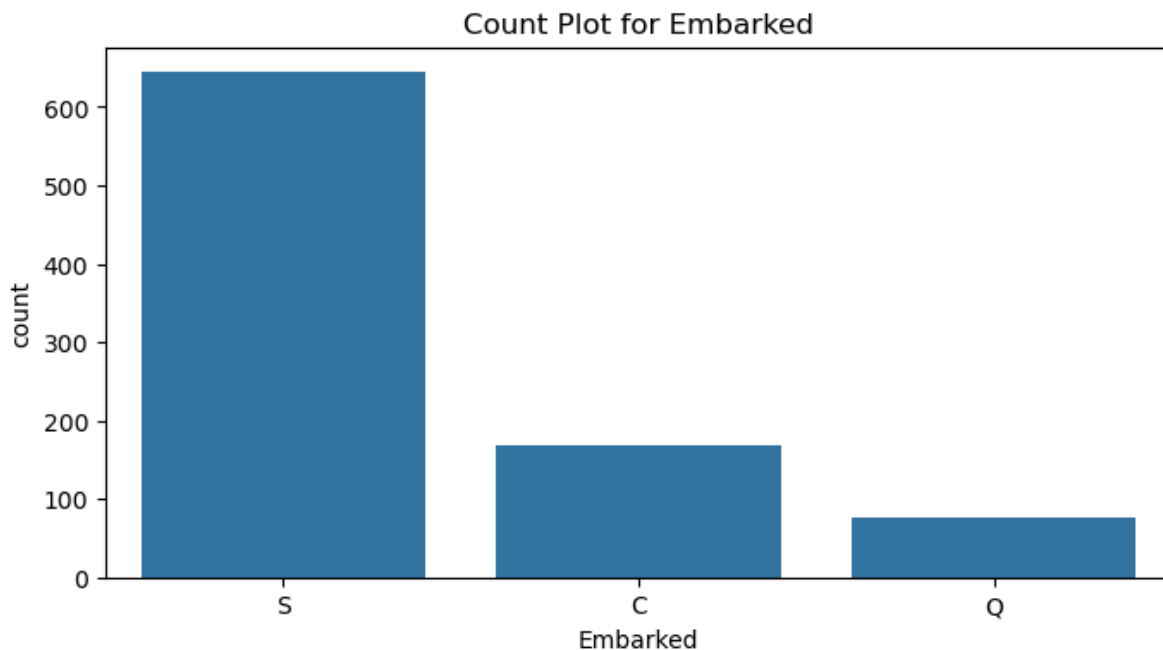
Histograms for Numeric Columns





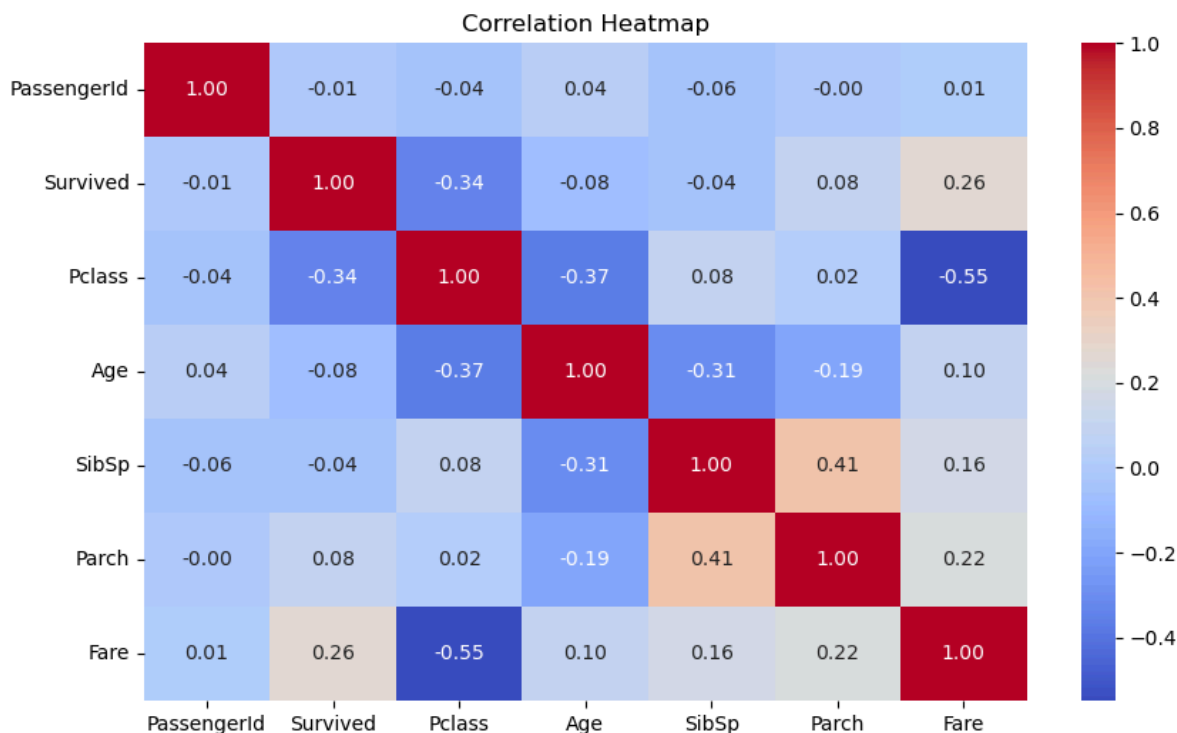
```
In [32]: # Plot bar charts for categorical columns
categorical_cols = ['Pclass', 'Sex', 'Embarked']
for col in categorical_cols:
    plt.figure(figsize=(8, 4))
    sns.countplot(data=df, x=col)
    plt.title(f"Count Plot for {col}")
    plt.show()
```



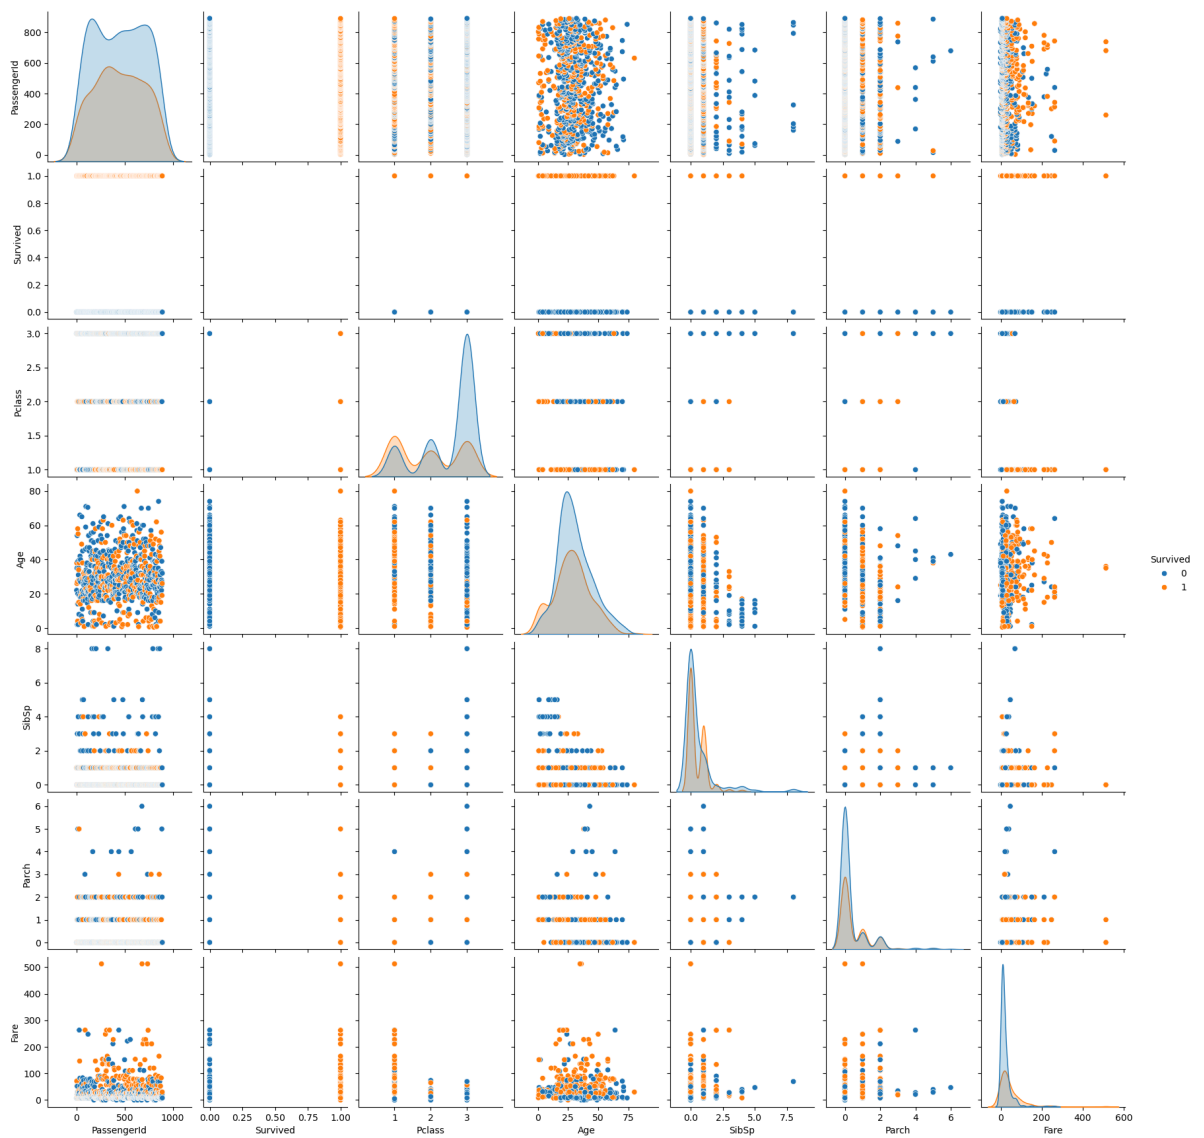


### 3- Correlation analysis

```
In [34]: # Correlation heatmap
tit_num = df.select_dtypes(include = ['float64', 'int64'])
plt.figure(figsize=(10, 6))
sns.heatmap(tit_num.corr(), annot=True, cmap='coolwarm', fmt=".2f")
plt.title("Correlation Heatmap")
plt.show()
```



```
In [46]: sns.pairplot(df, hue='Survived', vars=numeric_cols.columns)
plt.show()
```




```
In [44]: numeric_cols = df.select_dtypes(include=['float64', 'int64'])
correlation_matrix = numeric_cols.corr()
survived_corr = correlation_matrix['Survived'].sort_values(ascending=False)
print(survived_corr)
```

```
Survived    1.000000
Fare         0.257307
Parch        0.081629
PassengerId -0.005007
SibSp        -0.035322
Age          -0.077221
Pclass      -0.338481
Name: Survived, dtype: float64
```

## 4- Log Transformation


```
In [6]: sp_untransformed = sns.displot(df['Fare'])
```



```
In [7]: print("Skewness: %f" % df['Fare'].skew())
```

Skewness: 4.787317

```
In [10]: log_transformed = np.log(df['Fare'])  
sp_transformed = sns.displot(log_transformed)
```



```
In [13]: df['Fare'] = log_transformed
```

## 4- Handling Duplicates

```
In [14]: duplicate = df[df.duplicated(['PassengerId'])]
duplicate
```


```
Out[14]: PassengerId  Survived  Pclass  Name  Sex  Age  SibSp  Parch  Ticket  Fare  Cabin  Embarked
```

## 5- Missing Values

```
In [20]: total = df.isnull().sum().sort_values(ascending=False)
total_select = total.head(20)
total_select.plot(kind="bar", figsize = (8,6), fontsize = 10)

plt.xlabel("Columns", fontsize = 20)
plt.ylabel("Count", fontsize = 20)
plt.title("Total Missing Values", fontsize = 20)
```

```
Out[20]: Text(0.5, 1.0, 'Total Missing Values')
```



```
In [18]: mean = df["Age"].mean()
         mean
```

```
Out[18]: 29.69911764705882
```

```
In [19]: df["Age"].fillna(mean, inplace = True)
```

/tmp/ipykernel\_2378236/1730557404.py:1: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or 'df[col] = df[col].method(value)' instead, to perform the operation inplace on the original object.

```
df["Age"].fillna(mean, inplace = True)
```


```
In [21]: df["Cabin"].fillna(0, inplace = True)
```

```
In [39]: df["Embarked"].fillna('C', inplace = True)
```

```
In [24]: total = df.isnull().sum().sort_values(ascending=False)
         total_select = total.head(20)
         total_select.plot(kind="bar", figsize = (8,6), fontsize = 10)
```

```
plt.xlabel("Columns", fontsize = 20)
plt.ylabel("Count", fontsize = 20)
plt.title("Total Missing Values", fontsize = 20)
```


Out[24]: Text(0.5, 1.0, 'Total Missing Values')



## 6- Outliers

In [25]: `sns.boxplot(x=df['Fare'])`

Out[25]: <Axes: xlabel='Fare'>



This is real data, the are paid is just bigger than he rest of passengers.

## 7- Scaling

```
In [30]: # Check for NaN values
print("NaN values in Fare column:", df['Fare'].isnull().sum())

# Check for infinite values
print("Infinite values in Fare column:", df['Fare'].isin([float('inf'), float('-inf')]).sum())

# Inspect large values
print("Maximum value in Fare column:", df['Fare'].max())
```

```
NaN values in Fare column: 0
Infinite values in Fare column: 15
Maximum value in Fare column: 6.238967387173661
```

```
In [31]: # Replace infinite values with a large finite value (e.g., max fare)
df['Fare'].replace([float('inf'), float('-inf')], df['Fare'].max(), inplace=True)
```

/tmp/ipykernel\_2378236/4193003221.py:2: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or 'df[col] = df[col].method(value)' instead, to perform the operation inplace on the original object.

```
df['Fare'].replace([float('inf'), float('-inf')], df['Fare'].max(), inplace=True)
```



```
In [32]: from sklearn.preprocessing import MinMaxScaler, StandardScaler

# Initialize scalers
min_max_scaler = MinMaxScaler() # Scales data to range [0, 1]
standard_scaler = StandardScaler() # Scales data to mean 0 and standard deviation

# Apply scaling to numeric columns
numeric_cols = ['Age', 'Fare'] # Select numeric columns for scaling
df['Age_MinMax'] = min_max_scaler.fit_transform(df[['Age']])
df['Fare_Standard'] = standard_scaler.fit_transform(df[['Fare']])

# Preview scaled values
print(df[['Age', 'Age_MinMax', 'Fare', 'Fare_Standard']].head())
```

	Age	Age_MinMax	Fare	Fare_Standard
0	22.0	0.271174	1.981001	-0.997611
1	38.0	0.472229	4.266662	1.242540
2	26.0	0.321438	2.070022	-0.910362
3	35.0	0.434531	3.972177	0.953918
4	35.0	0.434531	2.085672	-0.895024

## 8- New variable creation

(a) Family Size Combine SibSp (siblings/spouses aboard) and Parch (parents/children aboard) to calculate the total family size:

```
In [33]: df['FamilySize'] = df['SibSp'] + df['Parch'] + 1 # Add 1 to include the passenger
print(df[['SibSp', 'Parch', 'FamilySize']].head())
```

	SibSp	Parch	FamilySize
0	1	0	2
1	1	0	2
2	0	0	1
3	1	0	2
4	0	0	1

(b) Title Extraction Extract titles from passenger names to analyze their influence on survival:

```
In [34]: df['Title'] = df['Name'].str.extract(' ([A-Za-z]+\.)', expand=False)
print(df['Title'].value_counts())
```

```
Title
Mr      517
Miss    182
Mrs     125
Master   40
Dr        7
Rev        6
Col         2
Mlle        2
Major        2
Ms           1
Mme          1
Don          1
Lady         1
Sir          1
Capt        1
Countess     1
Jonkheer     1
Name: count, dtype: int64
```

(c) Fare Per Person Calculate the fare per family member:

```
In [35]: df['FarePerPerson'] = df['Fare'] / df['FamilySize']
print(df[['Fare', 'FamilySize', 'FarePerPerson']].head())
```

	Fare	FamilySize	FarePerPerson
0	1.981001	2	0.990501
1	4.266662	2	2.133331
2	2.070022	1	2.070022
3	3.972177	2	1.986088
4	2.085672	1	2.085672

## 9- Encoding Categorical Features

```
In [40]: df['Sex_Encoded'] = df['Sex'].map({'male': 0, 'female': 1})
print(df[['Sex', 'Sex_Encoded']].head())
```

	Sex	Sex_Encoded
0	male	0
1	female	1
2	female	1
3	female	1
4	male	0

```
In [42]: # Get a Pd.Series consisting of all the string categoricals
one_hot_encode_cols = df.dtypes[df.dtypes == object] # filtering by string categoricals
one_hot_encode_cols = one_hot_encode_cols.index.tolist() # list of categorical fields
df[one_hot_encode_cols].head().T
```

```
Out[42]:
```

	0	1	2	3	4
<b>Name</b>	Braund, Mr. Owen Harris	Cumings, Mrs. John Bradley (Florence Briggs Th...	Heikkinen, Miss. Laina	Futrelle, Mrs. Jacques Heath (Lily May Peel)	Allen, Mr. William Henry
<b>Sex</b>	male	female	female	female	male
<b>Ticket</b>	A/5 21171	PC 17599	STON/O2. 3101282	113803	373450
<b>Cabin</b>	0	C85	0	C123	0
<b>Embarked</b>	S	C	S	S	S
<b>Title</b>	Mr	Mrs	Miss	Mrs	Mr

```
In [43]: # One-hot encode
df = pd.get_dummies(df, columns=['Embarked', 'Title'], drop_first=True)

# Preview encoded dataframe
print(df.head())
```

	PassengerId	Survived	Pclass	\
0	1	0	3	
1	2	1	1	
2	3	1	3	
3	4	1	1	
4	5	0	3	

	Name	Sex	Age	SibSp	\
0	Braund, Mr. Owen Harris	male	22.0	1	
1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	
2	Heikkinen, Miss. Laina	female	26.0	0	
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	
4	Allen, Mr. William Henry	male	35.0	0	

	Parch	Ticket	Fare	... Title_Major	Title_Master	\
0	0	A/5 21171	1.981001	... False	False	
1	0	PC 17599	4.266662	... False	False	
2	0	STON/O2. 3101282	2.070022	... False	False	
3	0	113803	3.972177	... False	False	
4	0	373450	2.085672	... False	False	

	Title_Miss	Title_Mlle	Title_Mme	Title_Mr	Title_Mrs	Title_Ms	\
0	False	False	False	True	False	False	
1	False	False	False	False	True	False	
2	True	False	False	False	False	False	
3	False	False	False	False	True	False	
4	False	False	False	True	False	False	

	Title_Rev	Title_Sir
0	False	False
1	False	False
2	False	False
3	False	False
4	False	False

[5 rows x 35 columns]

In [44]: `df.info()`

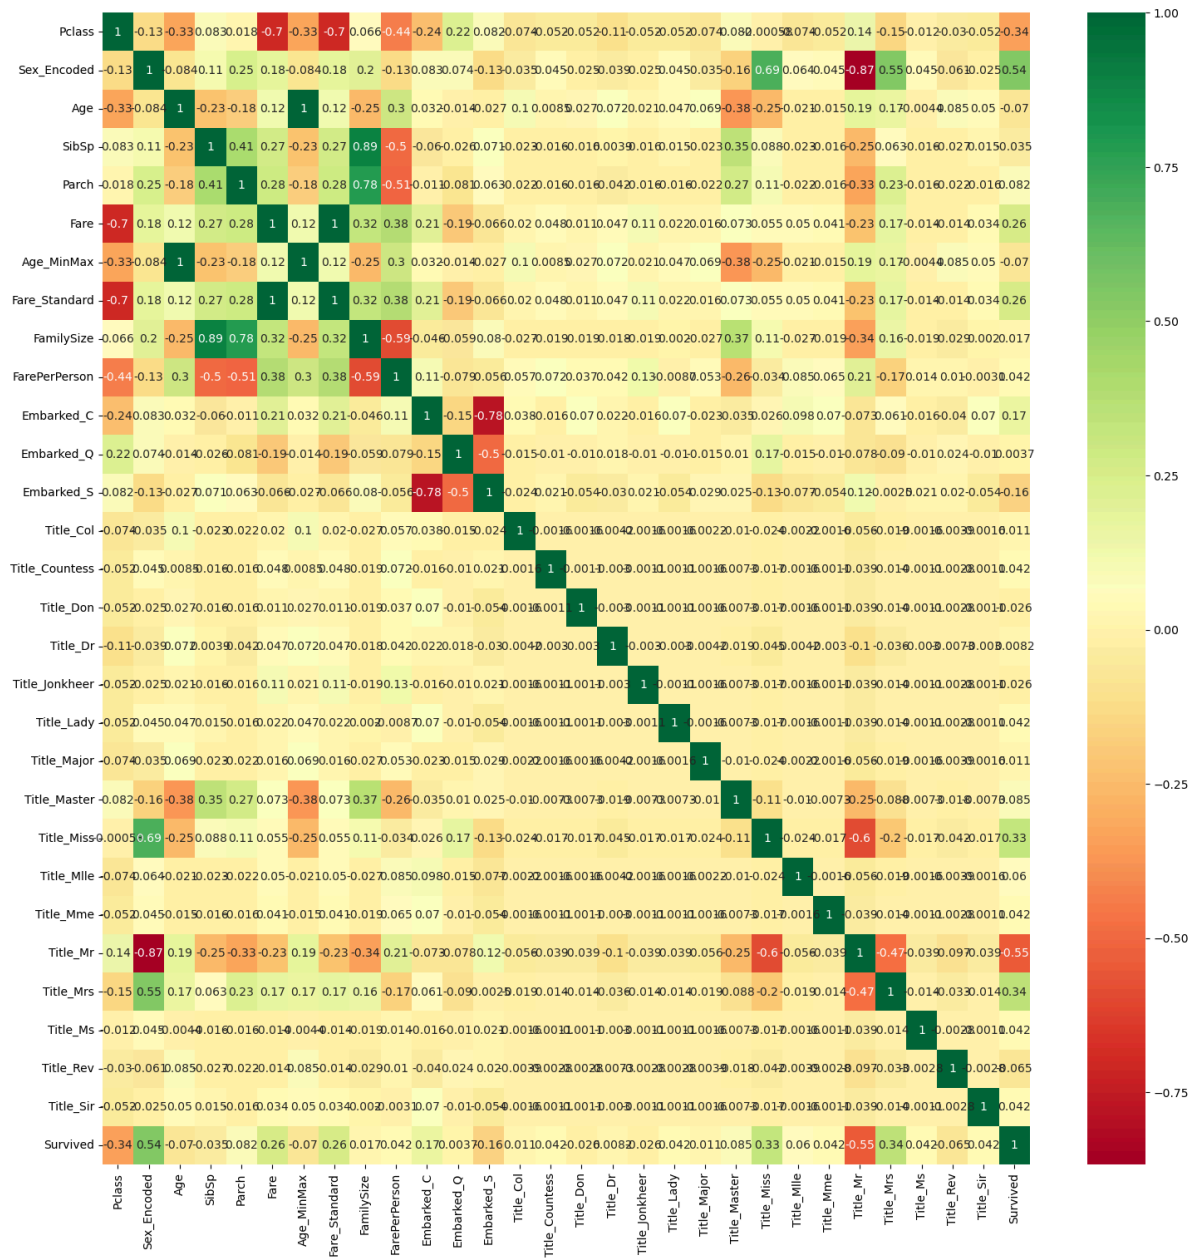
```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 35 columns):
#   Column                Non-Null Count  Dtype
---  -
0   PassengerId           891 non-null    int64
1   Survived              891 non-null    int64
2   Pclass                891 non-null    int64
3   Name                  891 non-null    object
4   Sex                   891 non-null    object
5   Age                   891 non-null    float64
6   SibSp                 891 non-null    int64
7   Parch                 891 non-null    int64
8   Ticket                891 non-null    object
9   Fare                  891 non-null    float64
10  Cabin                 891 non-null    object
11  Age_MinMax            891 non-null    float64
12  Fare_Standard         891 non-null    float64
13  FamilySize            891 non-null    int64
14  FarePerPerson         891 non-null    float64
15  Sex_Encoded           891 non-null    int64
16  Embarked_C            891 non-null    bool
17  Embarked_Q            891 non-null    bool
18  Embarked_S            891 non-null    bool
19  Title_Col             891 non-null    bool
20  Title_Countess        891 non-null    bool
21  Title_Don             891 non-null    bool
22  Title_Dr              891 non-null    bool
23  Title_Jonkheer        891 non-null    bool
24  Title_Lady            891 non-null    bool
25  Title_Major           891 non-null    bool
26  Title_Master          891 non-null    bool
27  Title_Miss            891 non-null    bool
28  Title_Mlle            891 non-null    bool
29  Title_Mme             891 non-null    bool
30  Title_Mr              891 non-null    bool
31  Title_Mrs             891 non-null    bool
32  Title_Ms              891 non-null    bool
33  Title_Rev             891 non-null    bool
34  Title_Sir             891 non-null    bool
dtypes: bool(19), float64(5), int64(7), object(4)
memory usage: 128.0+ KB
```

## 10- Feature Selection


```
In [48]: titanic = df.loc[:,['Pclass', 'Sex_Encoded',
    'Age', 'SibSp', 'Parch',
    'Fare', 'Age_MinMax', 'Fare_Standard',
    'FamilySize', 'FarePerPerson', 'Embarked_C',
    'Embarked_Q', 'Embarked_S', 'Title_Col', 'Title_Countess',
    'Title_Don', 'Title_Dr', 'Title_Jonkheer', 'Title_Lady', 'Title_Major',
    'Title_Master', 'Title_Miss', 'Title_Mlle', 'Title_Mme', 'Title_Mr',
    'Title_Mrs', 'Title_Ms', 'Title_Rev', 'Title_Sir', 'Survived']]
```

```
In [49]: plt.figure(figsize=(18,18))
sns.heatmap(titanic.corr(),annot=True,cmap='RdYlGn')


plt.show()
```



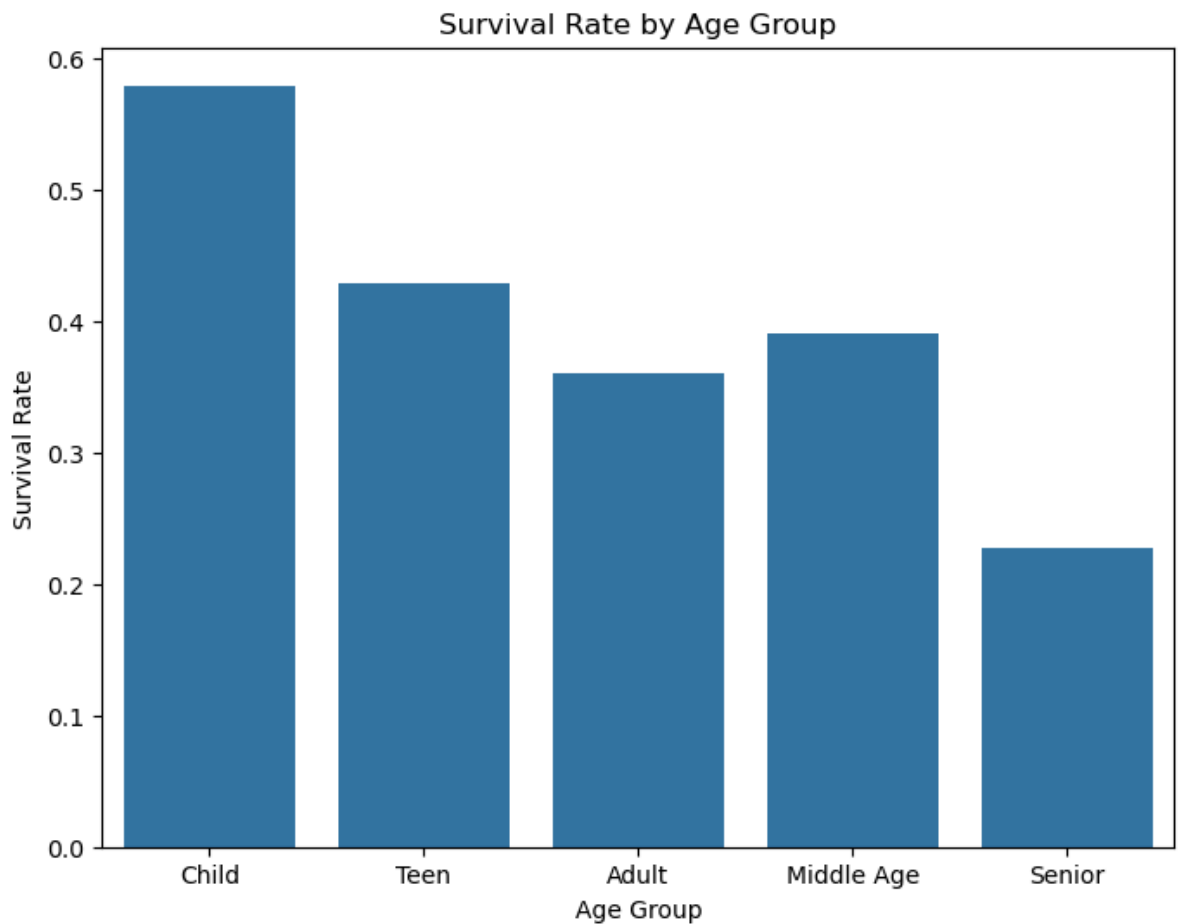
```
In [54]: # Survival rate by gender
plt.figure(figsize=(8, 6))
sns.barplot(data=df, x='Sex', y='Survived', errorbar=None)
plt.title("Survival Rate by Gender")
plt.ylabel("Survival Rate")
plt.xlabel("Gender")
plt.show()
```



```
In [53]: # Survival rate by passenger class
plt.figure(figsize=(8, 6))
sns.barplot(data=df, x='Pclass', y='Survived', errorbar=None)
plt.title("Survival Rate by Passenger Class")
plt.ylabel("Survival Rate")
plt.xlabel("Passenger Class")
plt.show()
```



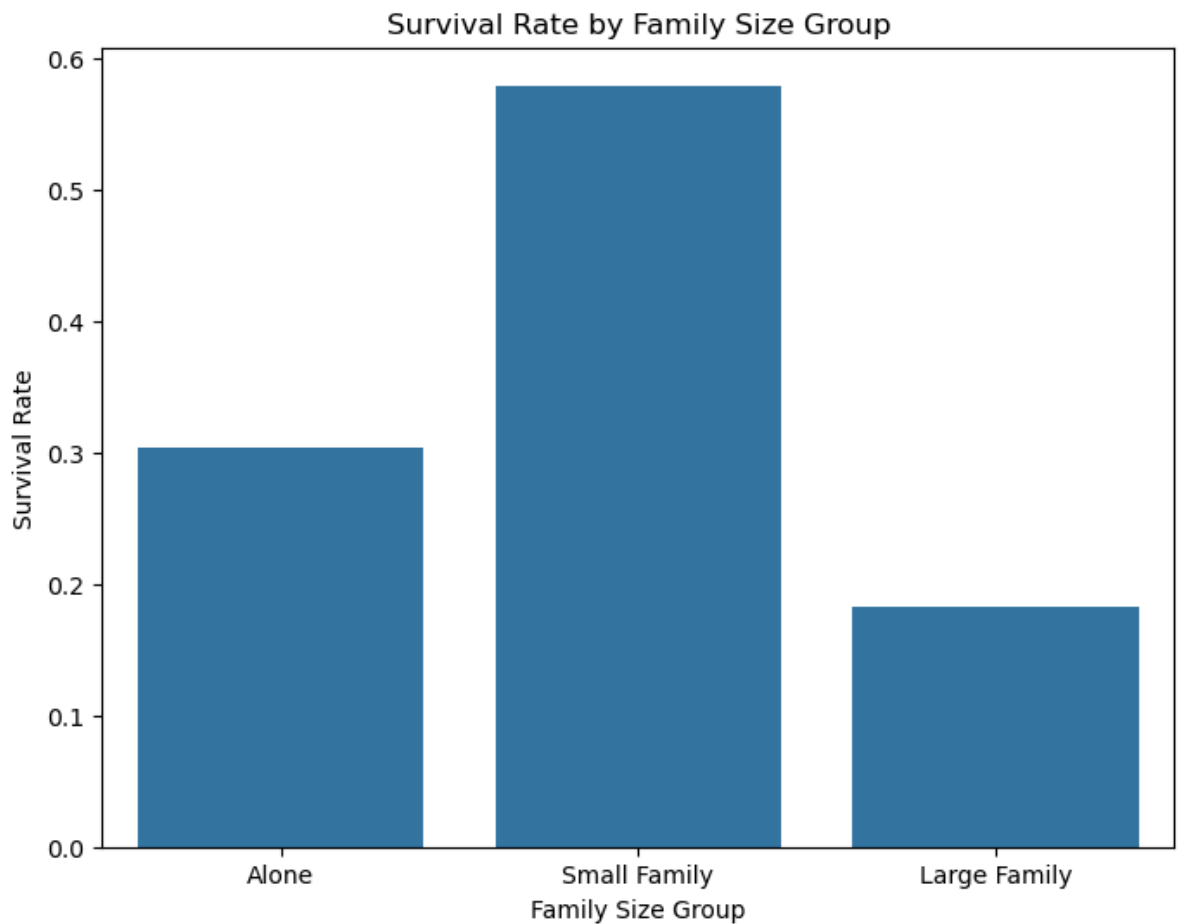
```
In [55]: # Bin ages for analysis
df['AgeGroup'] = pd.cut(df['Age'], bins=[0, 12, 18, 40, 60, 80], labels=['Child', '
# Survival rate by age group
plt.figure(figsize=(8, 6))
sns.barplot(data=df, x='AgeGroup', y='Survived', errorbar=None)
plt.title("Survival Rate by Age Group")
plt.ylabel("Survival Rate")
plt.xlabel("Age Group")
plt.show()
```



```
In [56]: # Bin family sizes for analysis
df['FamilySizeGroup'] = pd.cut(df['FamilySize'], bins=[0, 1, 4, 10], labels=['Alone', 'Small', 'Medium', 'Large'])

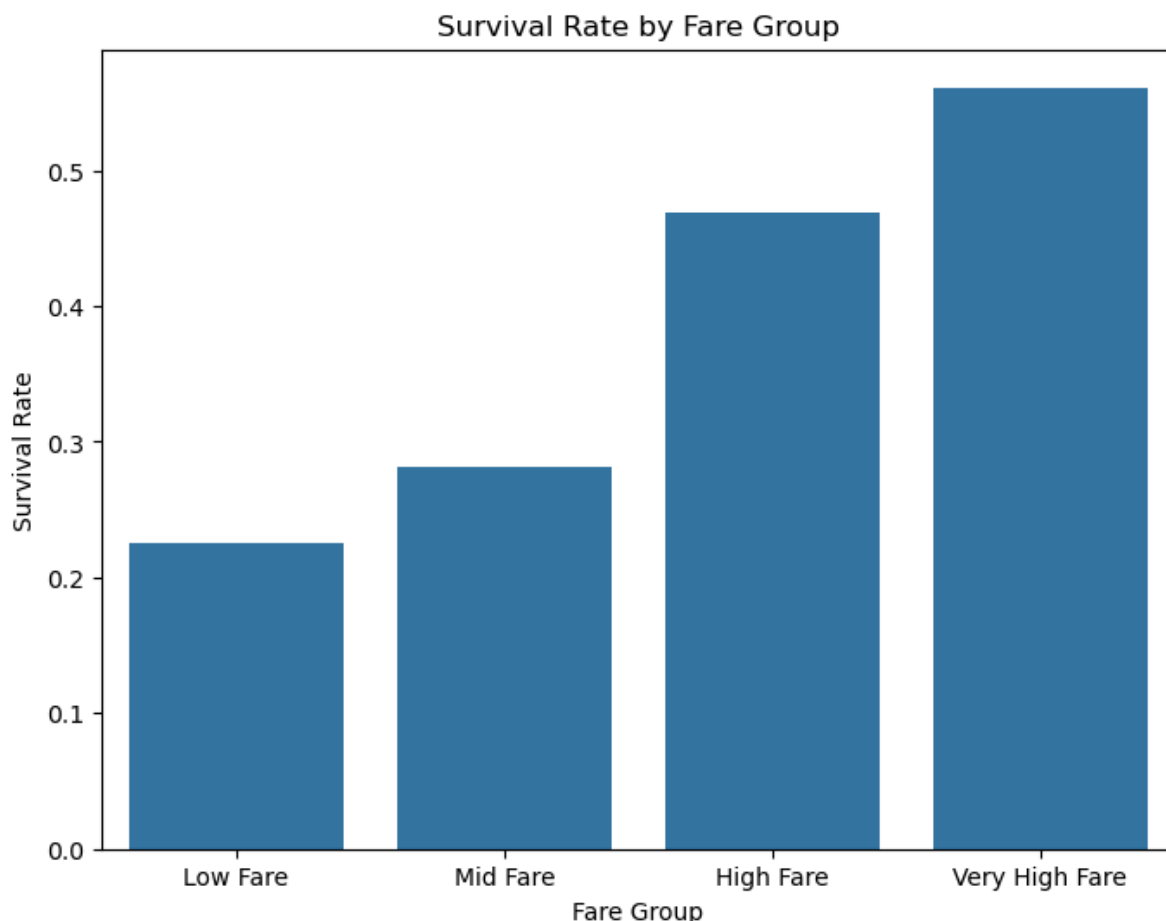
# Survival rate by family size group
plt.figure(figsize=(8, 6))
sns.barplot(data=df, x='FamilySizeGroup', y='Survived', errorbar=None)
plt.title("Survival Rate by Family Size Group")
plt.ylabel("Survival Rate")
plt.xlabel("Family Size Group")
plt.show()
```





```
In [57]: # Bin fares for analysis
df['FareGroup'] = pd.qcut(df['Fare'], q=4, labels=['Low Fare', 'Mid Fare', 'High Fa

# Survival rate by fare group
plt.figure(figsize=(8, 6))
sns.barplot(data=df, x='FareGroup', y='Survived', errorbar=None)
plt.title("Survival Rate by Fare Group")
plt.ylabel("Survival Rate")
plt.xlabel("Fare Group")
plt.show()
```



## STEP 4: Key Findings and Insights, which synthesizes the results of Exploratory Data Analysis in an insightful and actionable manner

### Synthesis from Exploratory Data Analysis (EDA):

#### Correlation

After the whole EDA, grouping and encoding the variables, we can see that the most correlated features with the survival rate are: Passenger Class, Sex, Age, Family Size and Fare. The Embarked place also had a slightly influence. The Title is not that important as the gender (The titles applied to women had bigger survival rates than those applied to men)

#### Survival Rate by Gender:

Females had a significantly higher survival rate compared to males. Gender played a critical role in survival, likely due to the "women and children first" evacuation policy.

#### Survival Rate by Passenger Class:

First-class passengers had a much higher survival rate compared to second-class and third-class. Socioeconomic status significantly influenced survival chances.

#### Survival Rate by Age:

Children (ages 0–12) had a higher survival rate compared to adults. Younger passengers had a better chance of survival.

**Impact of Family Size:**

Passengers traveling alone (FamilySize = 1) had a lower survival rate than those with small families (FamilySize = 2–4). Traveling with family increased survival chances, but very large family groups faced lower survival rates.

**Impact of Fare:**

Passengers who paid higher fares generally had better survival chances. Fare serves as a proxy for class or cabin quality, further emphasizing the role of socioeconomic factors.

## STEP 5: Formulating 3 hypothesis about this data

**Hypothesis 1 (Gender):**

*Null Hypothesis ( $H_0$ ):* Survival is independent of the passenger's gender.

*Alternative Hypothesis ( $H_1$ ):* Survival is dependent on the passenger's gender.

**Hypothesis 2 (Passenger Class):**

*Null Hypothesis ( $H_0$ ):* Survival rates are the same across all passenger classes.

*Alternative Hypothesis ( $H_1$ ):* Survival rates vary across passenger classes.

**Hypothesis 3 (Age):**

*Null Hypothesis ( $H_0$ ):* Survival is independent of the passenger's age.

*Alternative Hypothesis ( $H_1$ ):* Younger passengers had higher survival rates.

## STEP 6: Conducting a formal significance test for one of the hypotheses and discuss the results

**Hypothesis 1 (Gender):**

*Null Hypothesis ( $H_0$ ):* Survival is independent of the passenger's gender.

*Alternative Hypothesis ( $H_1$ ):* Survival is dependent on the passenger's gender.

```
In [60]: from scipy.stats import chi2_contingency

# Create a contingency table for Gender and Survival
contingency_table = pd.crosstab(df['Sex'], df['Survived'])

# Perform the Chi-Square Test
chi2, p, dof, expected = chi2_contingency(contingency_table)

# Results
print("Chi-Square Statistic:", chi2)
print("p-value:", p)

# Interpretation
if p < 0.05:
```

```
print("Reject the null hypothesis: Gender and Survival are dependent.")  
else:  
    print("Fail to reject the null hypothesis: Gender and Survival are independent.")
```

Chi-Square Statistic: 260.71702016732104

p-value: 1.1973570627755645e-58

Reject the null hypothesis: Gender and Survival are dependent.

## STEP 7: Suggestions for next steps in analyzing this data

### Suggestions for Next Steps in Analyzing This Data

#### **Incorporate Interaction Effects:**

Study how the interaction between Pclass and Gender influences survival rates.

#### **Build Predictive Models:**

Train a Machine Learning model to predict survival, using features like Pclass, Sex, Age, and Fare.

#### **Explore External Data Sources:**

Merge with external datasets, such as ship manifests or historical records, to enrich the analysis.

## STEP 8: A paragraph that summarizes the quality of this data set and a request for additional data if needed

The dataset contains critical variables for survival prediction (e.g., Sex, Pclass, Fare). Data is relatively clean and well-documented.

However, missing values in Age (~20%) and Cabin (~77%) limit analysis and predictive accuracy. Lack of detailed data about passengers' health or physical conditions.

#### *Request for Additional Data:*

- Passenger Details: Information about passengers' health, physical conditions, or mobility status.
- Evacuation Process: Data on lifeboat allocation and boarding sequences.
- Crew Details: Adding data on the crew could provide insights into survival rates among different groups on the ship.